


EXAM PRESENTATION





01

INTRODUCTION

02

DATA
ANALYSIS

03

DATA
PREPROCESSING

04

DATA
PROCESSING



Raw data

age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex	capital.gain	capital.loss	hours.per.week	native.country	income
90	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Female	0	4356	40	United-States	<=50K
82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Female	0	4356	18	United-States	<=50K
66	?	186061	Some-college	10	Widowed	?	Unmarried	Black	Female	0	4356	40	United-States	<=50K
54	Private	140359	7th-8th	4	Divorced	Machine-operatives	Unmarried	White	Female	0	3900	40	United-States	<=50K
41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Female	0	3900	40	United-States	<=50K
34	Private	216864	HS-grad	9	Divorced	Other-service	Unmarried	White	Female	0	3770	45	United-States	<=50K
38	Private	150601	10th	6	Separated	Adm-clerical	Unmarried	White	Male	0	3770	40	United-States	<=50K
74	State-gov	88638	Doctorate	16	Never-married	Prof-specialty	Other-relative	White	Female	0	3683	20	United-States	>50K
68	Federal-gov	422013	HS-grad	9	Divorced	Prof-specialty	Not-in-family	White	Female	0	3683	40	United-States	<=50K
41	Private	70037	Some-college	10	Never-married	Craft-repair	Unmarried	White	Male	0	3004	60	?	>50K
45	Private	172274	Doctorate	16	Divorced	Prof-specialty	Unmarried	Black	Female	0	3004	35	United-States	>50K
38	Self-employed	164526	Prof-school	15	Never-married	Prof-specialty	Not-in-family	White	Male	0	2824	45	United-States	>50K
52	Private	129177	Bachelors	13	Widowed	Other-service	Not-in-family	White	Female	0	2824	20	United-States	>50K
32	Private	136204	Masters	14	Separated	Exec-managerial	Not-in-family	White	Male	0	2824	55	United-States	>50K
51	?	172175	Doctorate	16	Never-married	?	Not-in-family	White	Male	0	2824	40	United-States	>50K
46	Private	45363	Prof-school	15	Divorced	Prof-specialty	Not-in-family	White	Male	0	2824	40	United-States	>50K
45	Private	172822	11th	7	Divorced	Transportation	Not-in-family	White	Male	0	2824	76	United-States	>50K
57	Private	317847	Masters	14	Divorced	Exec-managerial	Not-in-family	White	Male	0	2824	50	United-States	>50K
22	Private	119592	Assoc-acad	12	Never-married	Handlers-cleaners	Not-in-family	Black	Male	0	2824	40	?	>50K
34	Private	203034	Bachelors	13	Separated	Sales	Not-in-family	White	Male	0	2824	50	United-States	>50K
37	Private	188774	Bachelors	13	Never-married	Exec-managerial	Not-in-family	White	Male	0	2824	40	United-States	>50K

Data analysis

General info


```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
 #   Column                Non-Null Count  Dtype  
---  --
 0   age                   32561 non-null  int64  
 1   workclass              32561 non-null  object  
 2   fnlwgt                 32561 non-null  int64  
 3   education              32561 non-null  object  
 4   education.num          32561 non-null  int64  
 5   marital.status         32561 non-null  object  
 6   occupation             32561 non-null  object  
 7   relationship           32561 non-null  object  
 8   race                   32561 non-null  object  
 9   sex                    32561 non-null  object  
10   capital.gain           32561 non-null  int64  
11   capital.loss           32561 non-null  int64  
12   hours.per.week         32561 non-null  int64  
13   native.country         32561 non-null  object  
14   income                 32561 non-null  object  
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

```
df.describe()
```


	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.week
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437456
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347429
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

Outliers Detection




'hours.per.week'

'education.num'





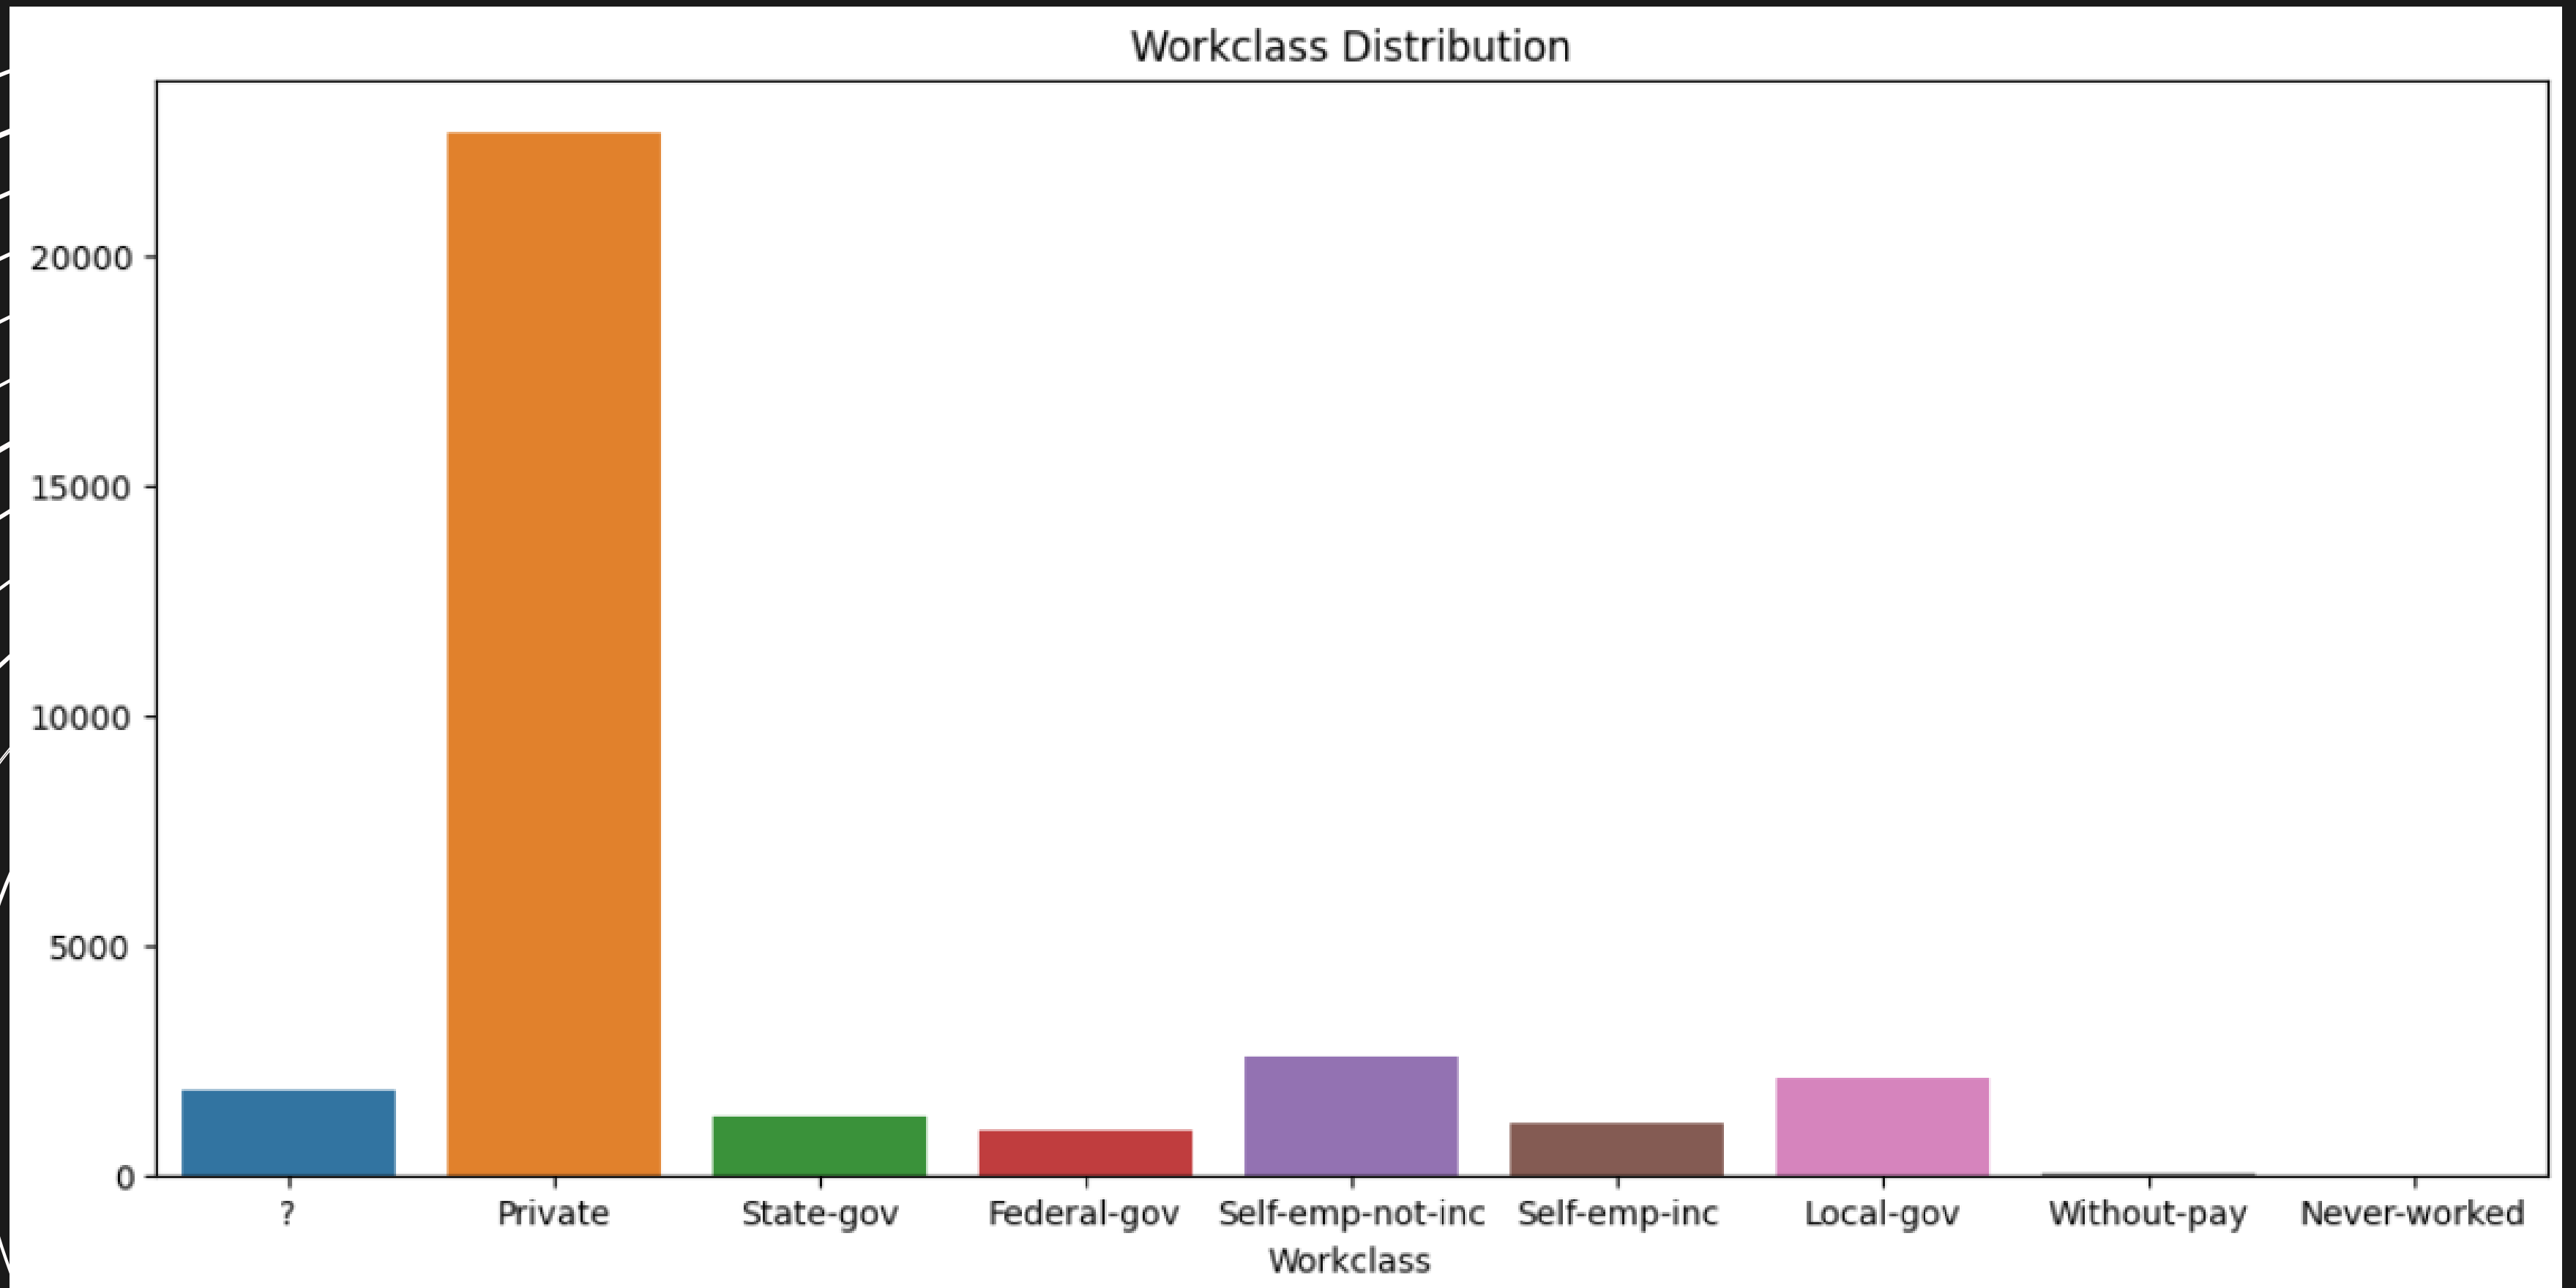
'age'

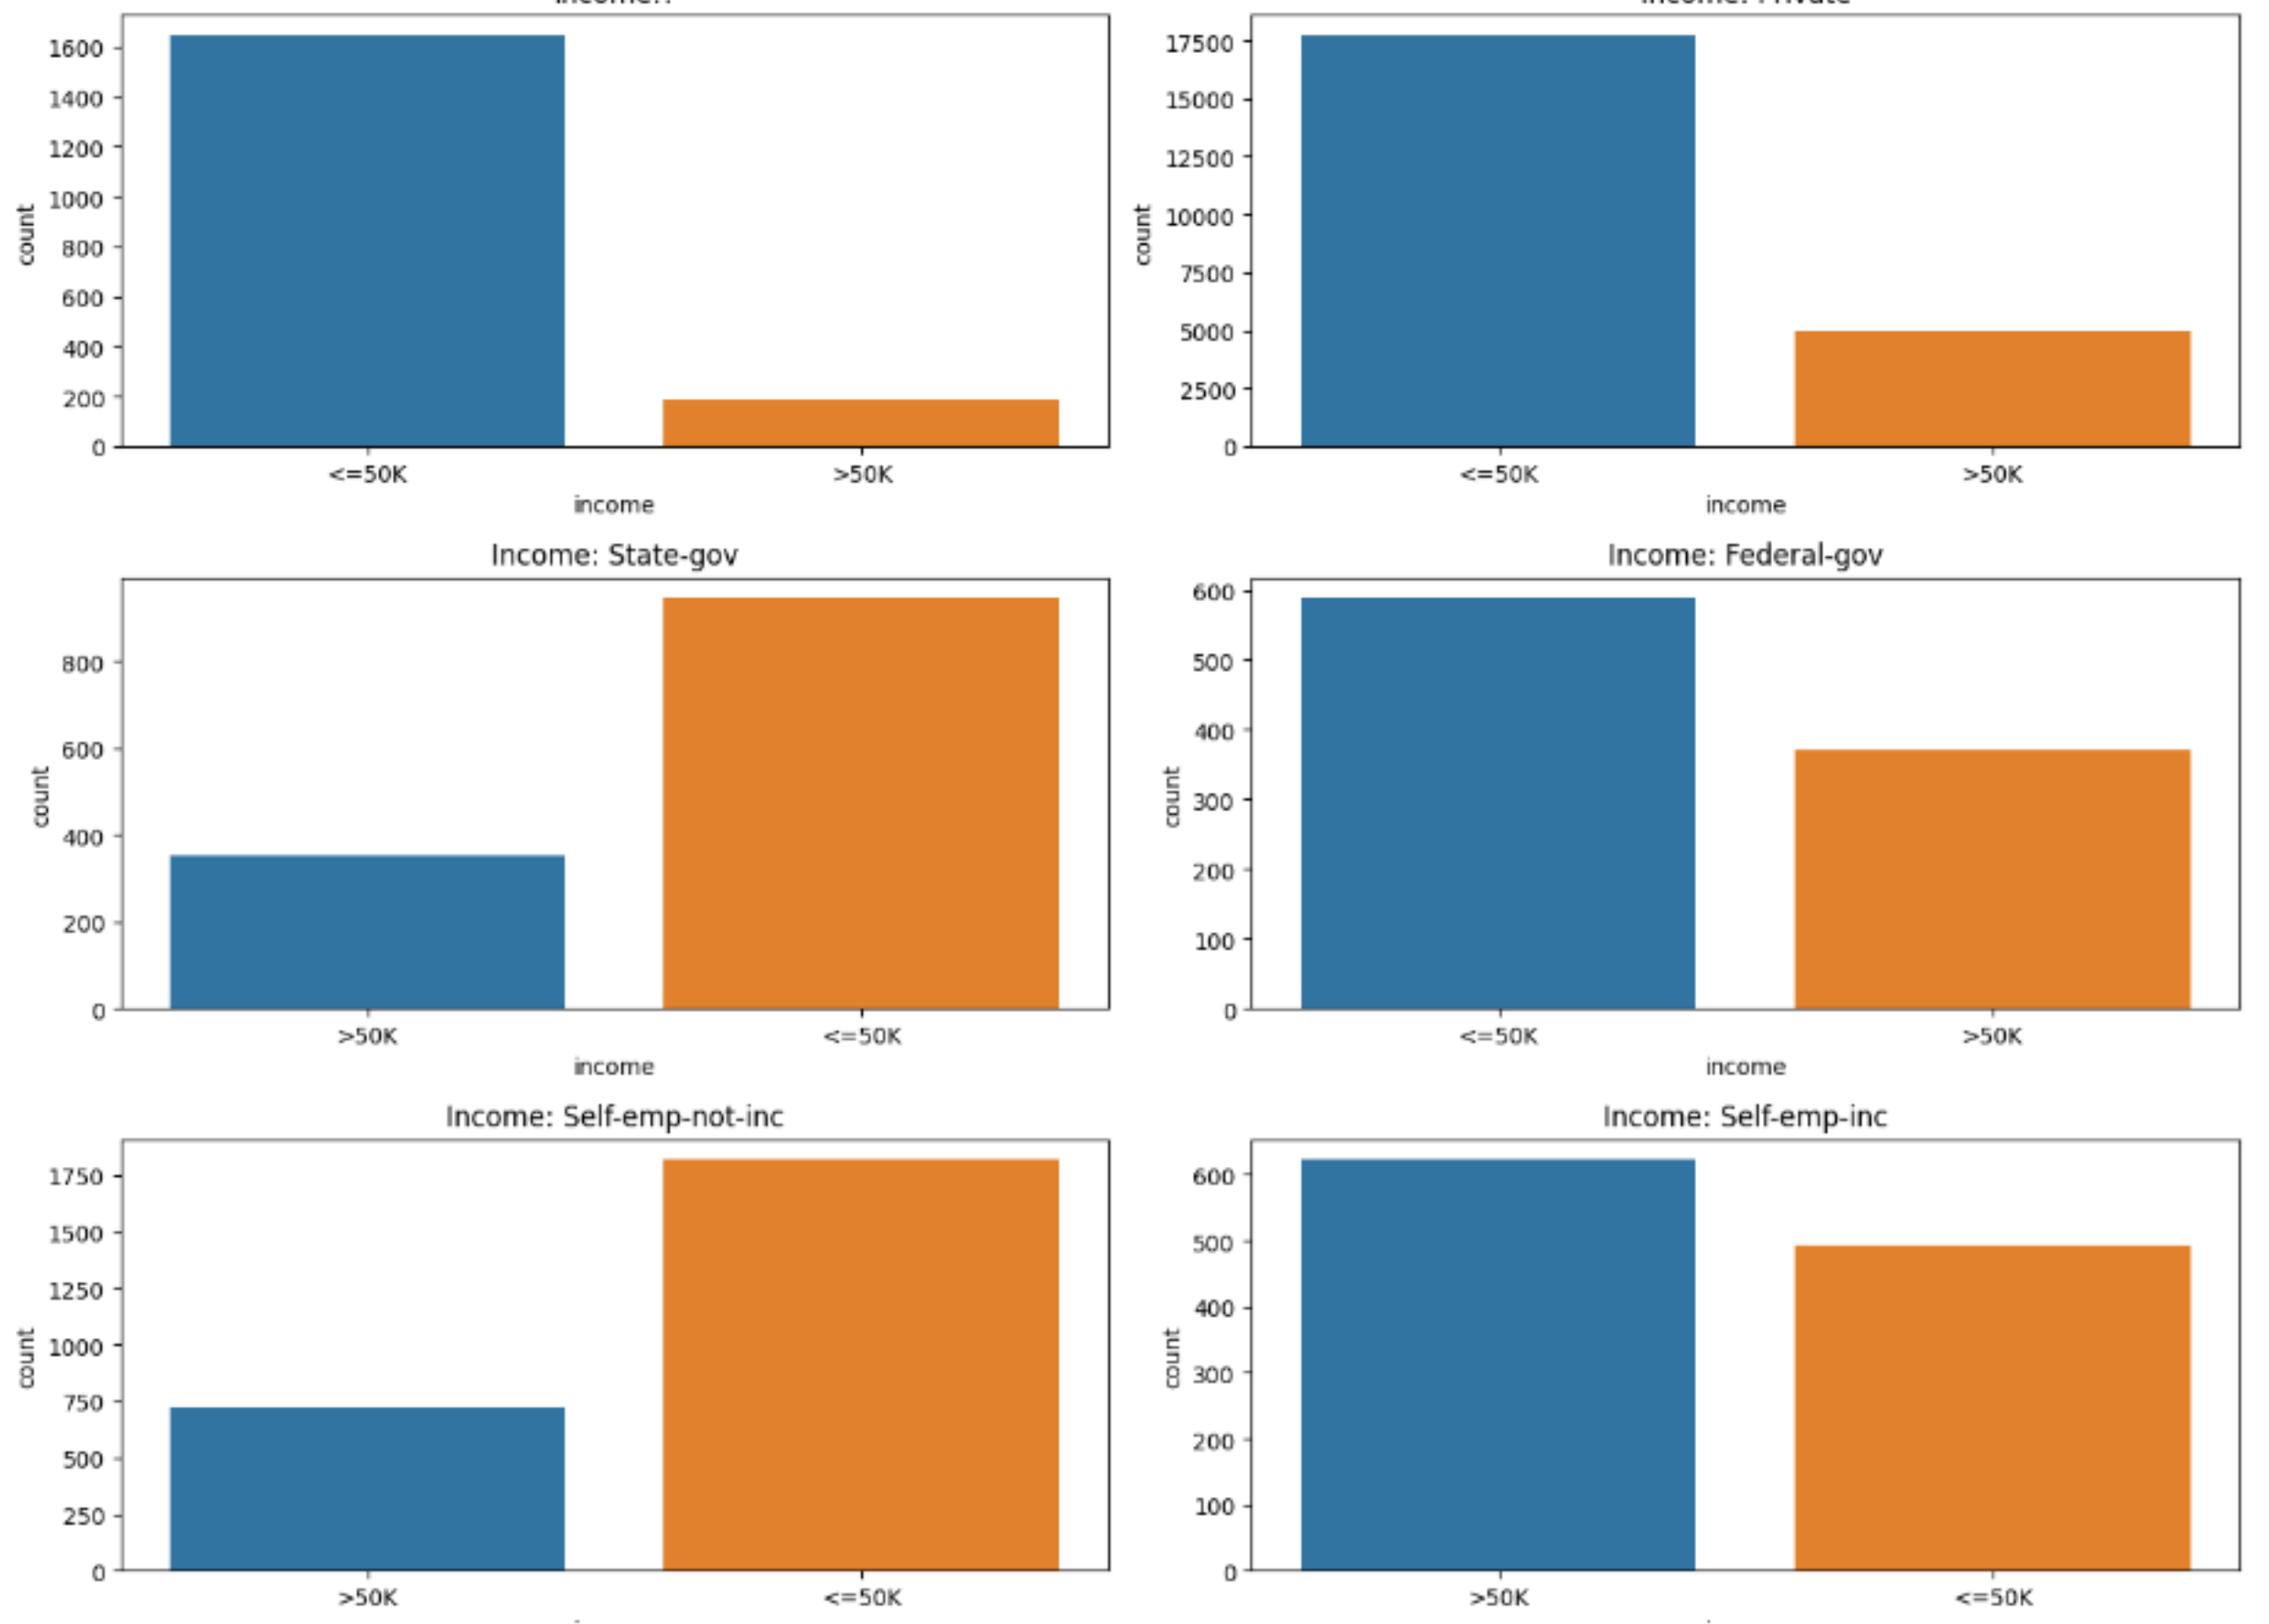


```
age - 0.43642622245443646
education.num - 3.6665949534376248
capital.gain - 8.335126164059378
capital.loss - 4.668531210621754
hours.per.week - 27.66696376432984
```

Interquartile range (IQR)


Workclass Distribution






We can see that people working at Federal-gov and Self-emp-inc have a higher income comparative to other workclasses, and the private workclass is the most popular one.

How does age influence capital gain and income?

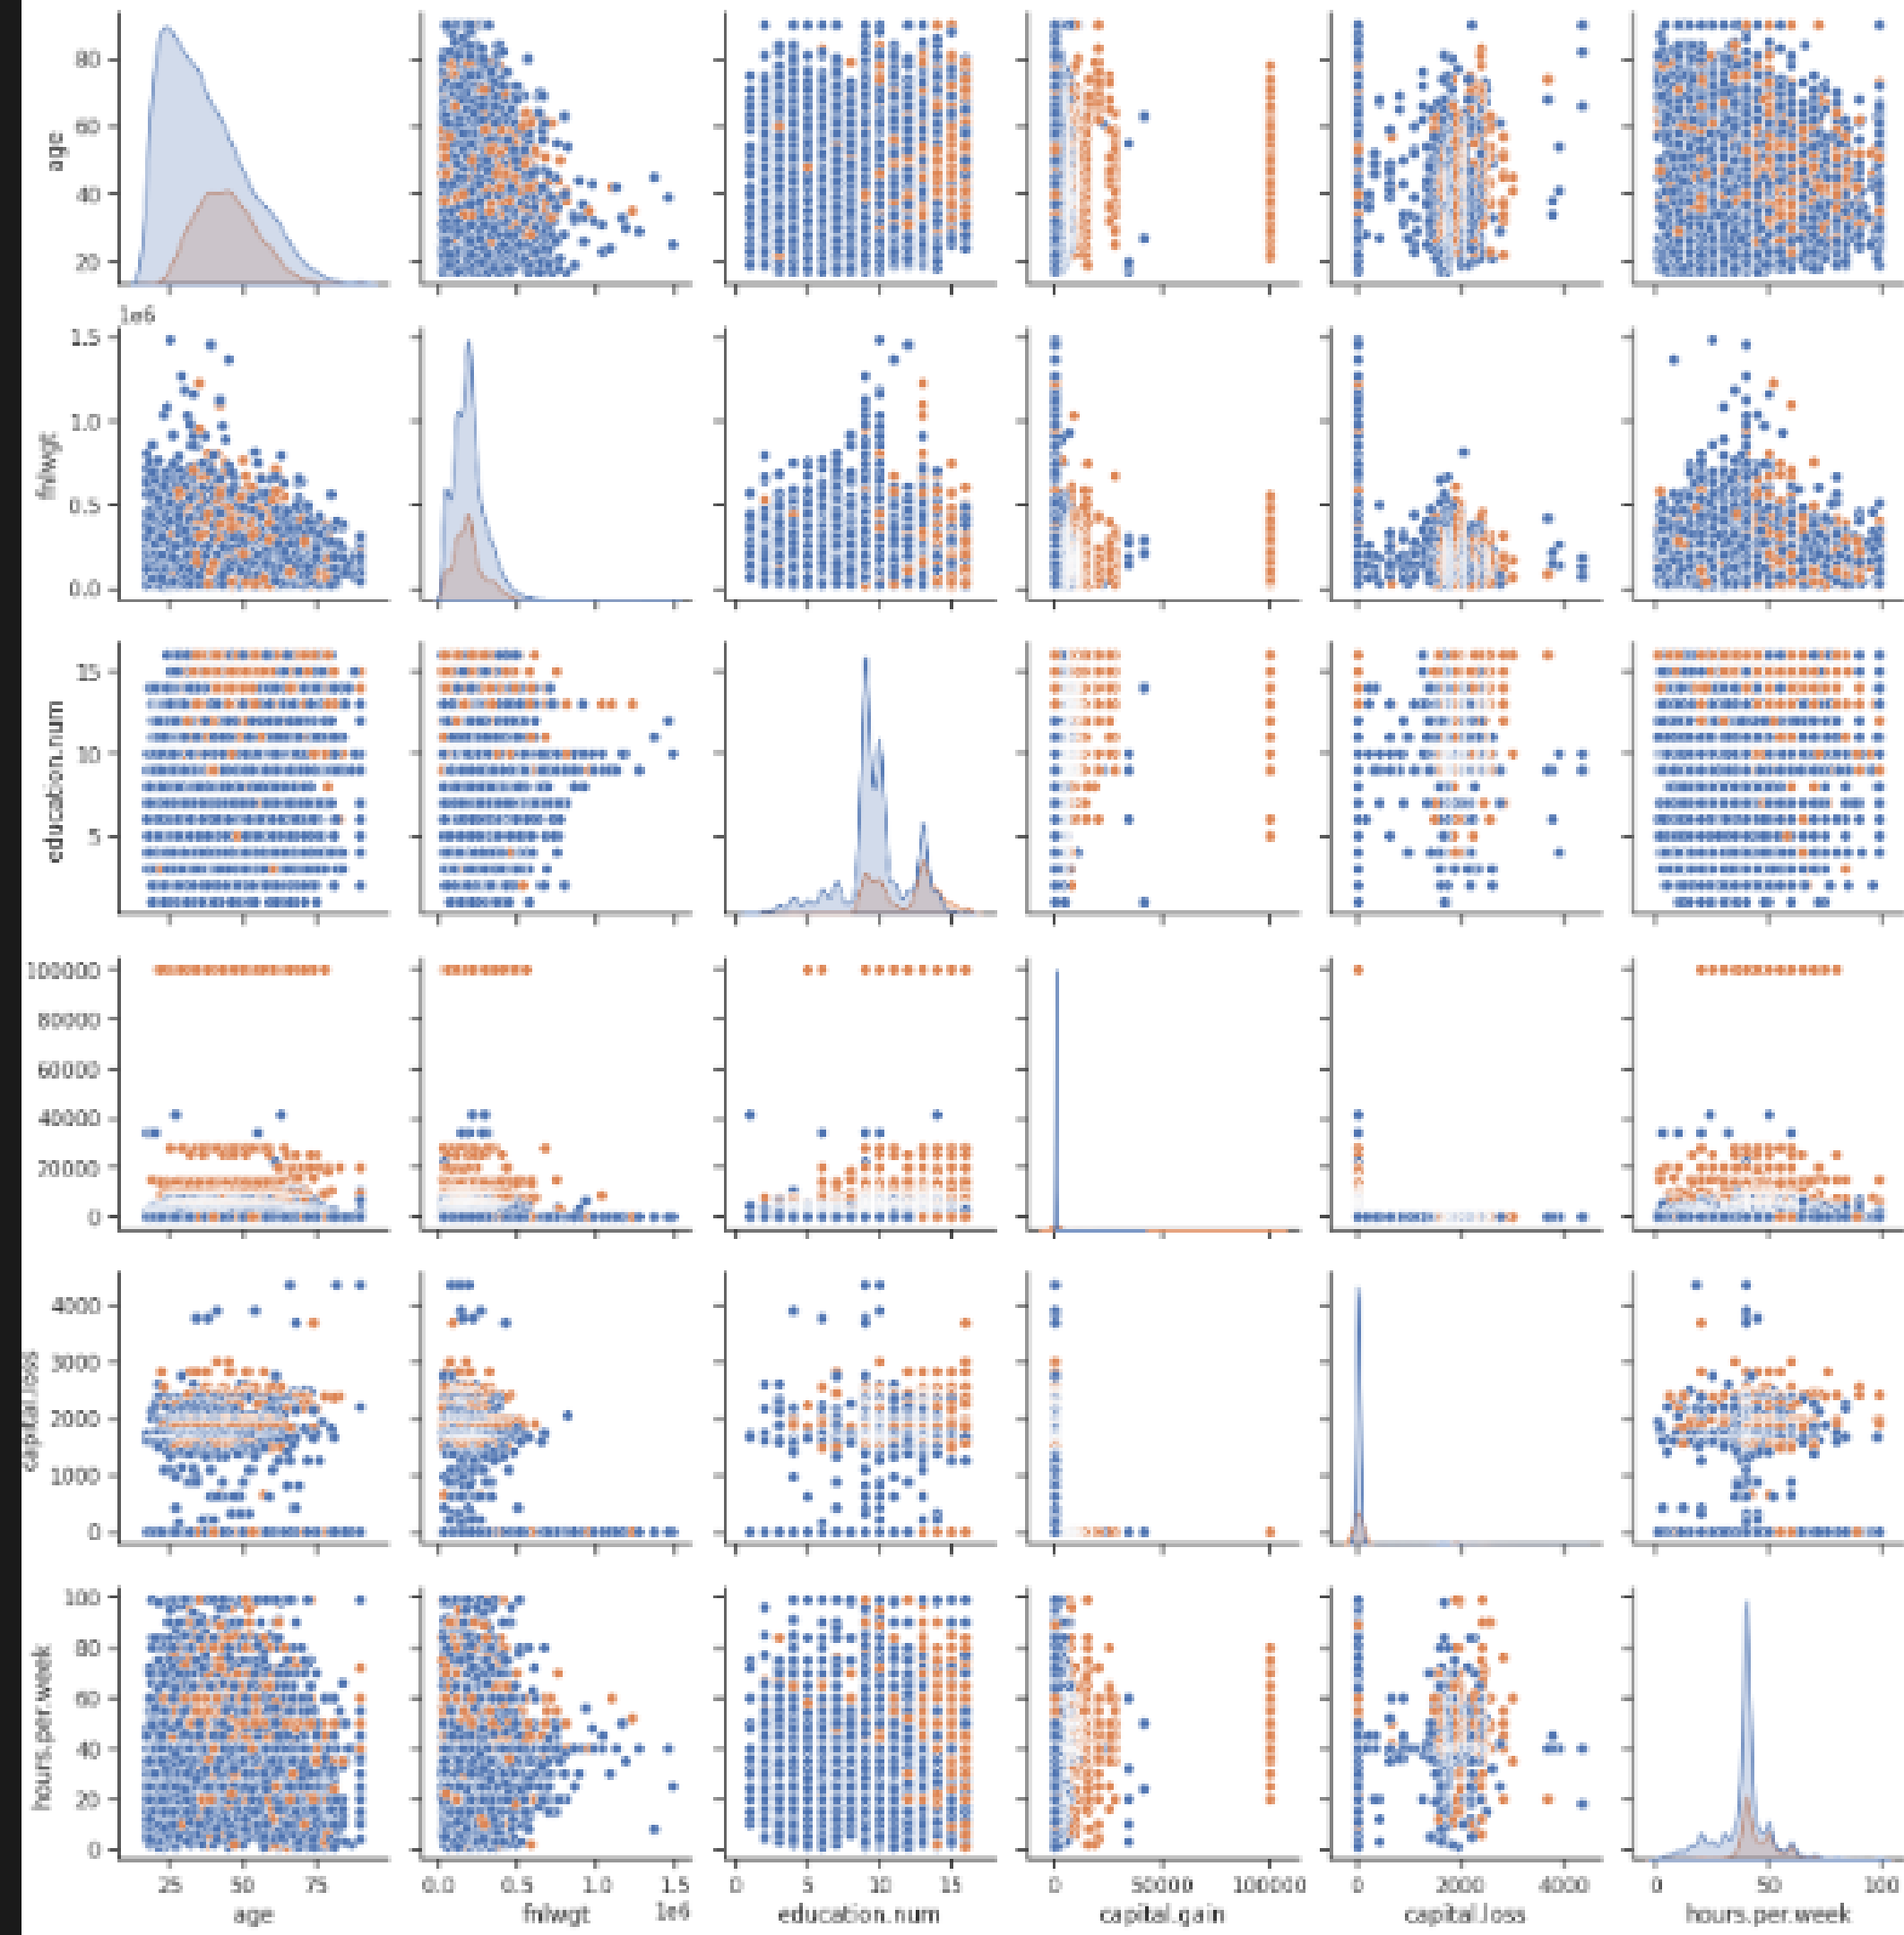


According to the plot above older people have a bigger income



Older people also reach a capital.gain but the values are not so high

We can notice that education.num has a linear relationship with income, the same relation can be noticed with hours.per.week and age columns



Data Engineering

Total number of '?' in df

```
k = 0
lista = len(df.columns)
for i in range(lista):
    for j in df[df.columns[i]]:
        if j == '?':
            k = k + 1
print(k)
```

4261

check the case when '?' is present in workclass and occupation

```
k = 0
for i in range(len(df)):
    if df.iloc[i,1]==df.iloc[i,6]=='?':
        k = k + 1
print(k)
```

1836

check the case when '?' is present in workclass and native.country

```
k = 0
for i in range(len(df)):
    if df.iloc[i,1]==df.iloc[i,14]=='?':
        k = k + 1
print(k)
```

0

check the case when '?' is present in native.country and occupation

```
k = 0
for i in range(len(df)):
    if df.iloc[i,6]==df.iloc[i,14]=='?':
        k = k + 1
print(k)
```

0

check the case when '?' is present in workclass , occupation and native.country

```
k = 0
for i in range(len(df)):
    if df.iloc[i,1]==df.iloc[i,14]==df.iloc[i,6]=='?':
        k = k + 1
print(k)
```

0

Analyzing the '?'

So far we can observe that the number of '?' in columns[occupation and workclass] is almost the same which means that we can not drop them since they have a valuable information [If X doesn't have a workclass results that it doesn't have an accupation]

```
def replace_inter(df):
    df['workclass'] = df['workclass'].str.replace('?', 'unemployed')
    df['occupation'] = df['occupation'].str.replace('?', 'no_occupation')
    df['native.country'] = df['native.country'].str.replace('?', 'unknown')
    return df
df = replace_inter(df)
```

```
k = 0
for i in df.workclass:
    if i == '?':
        k = k + 1
print(k)
```

1836

Check the number of '?' in occupation

```
k = 0
for i in df.occupation:
    if i == '?':
        k = k + 1
print(k)
```

1843

Check the number of '?' in native.country

```
k = 0
for i in df['native.country']:
    if i == '?':
        k = k + 1
print(k)
```

582

Categorical features

Binary_encoder

```
def binary_encoder(df, column_list):  
    for i in column_list:  
        rep = len(df[i].unique())  
        if rep > 3:  
            li1.append(i)  
    encoder = ce.BinaryEncoder(cols=li1, return_df = True)  
    df = encoder.fit_transform(df)  
    return df  
  
df = binary_encoder(df, lista)
```

Map function

```
def map_binar(df):  
    df['income'] = df['income'].map({'<=50K' : 0, '>50K' : 1})  
    df['sex'] = df['sex'].map({'Female' : 0, 'Male' : 1})  
    return df  
df = map_binar(df)
```


Class balancing

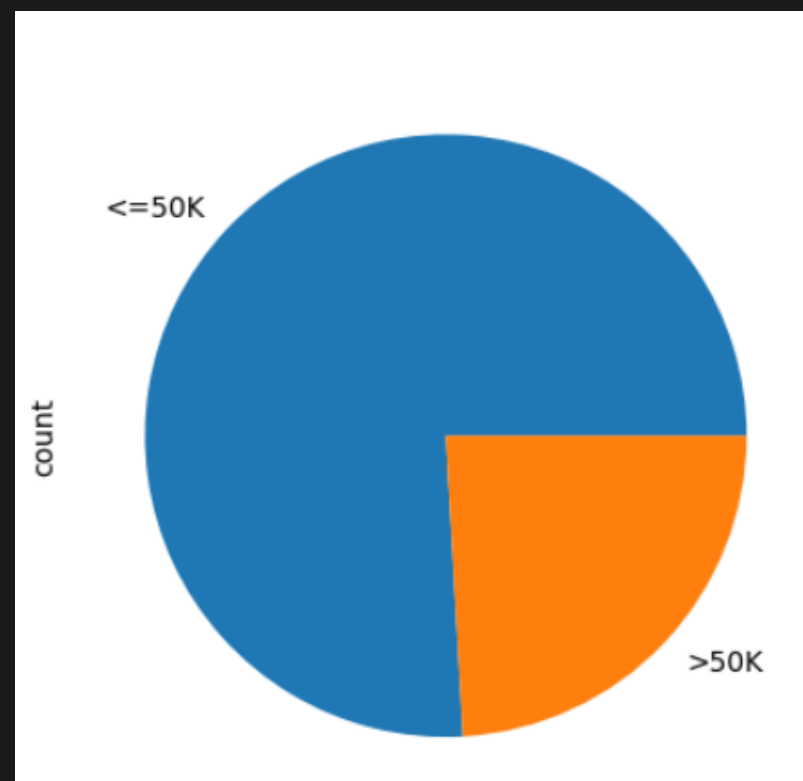
SMOTENC

```
from imblearn.over_sampling import SMOTENC
from sklearn.model_selection import train_test_split

def get_smotenc(df, target, num):
    X = df.drop(target, axis=1)
    Y = df[target]
    X_train, X_test, y_train, y_test = train_test_split(X, Y, random_state = 12, test_size=0.25)
    sm = SMOTENC(random_state=42, categorical_features=[num])
    X_res, y_res = sm.fit_resample(X_train, y_train)
    y_res.value_counts().plot.pie()
    df_balanced = pd.concat([X_res, y_res], axis = 1)
    return df_balanced
df = get_smotenc(df, 'income', 34)
df
```


Before

32537



After

37186





Outliers

IsolationForest
EllipticEnvelope
LocalOutlierFactor

*According to the result above
IsolationForest algorithm has
the best result since it is
capable to identify outliers in
non-linear data distributions.*

```
age - 0.47047492035267097
education.num - 0.048158850114840335
capital.gain - 13.2807290509002
capital.loss - 6.271764095724976
hours.per.week - 9.36504408387049
```

```
age - 0.7111482953357038
education.num - 0.40039441838228707
capital.gain - 12.872381749185765
capital.loss - 0.020916126333403055
hours.per.week - 16.89425404129441
```

```
age - 0.6837002328117172
education.num - 0.4449945480800401
capital.gain - 10.897945952317803
capital.loss - 5.160168567471193
hours.per.week - 12.775174608787907
```

Feature Selection

```
from kydavra import PValueSelector
from kydavra import BregmanDivergenceSelector
from kydavra import JensenShannonSelector
```

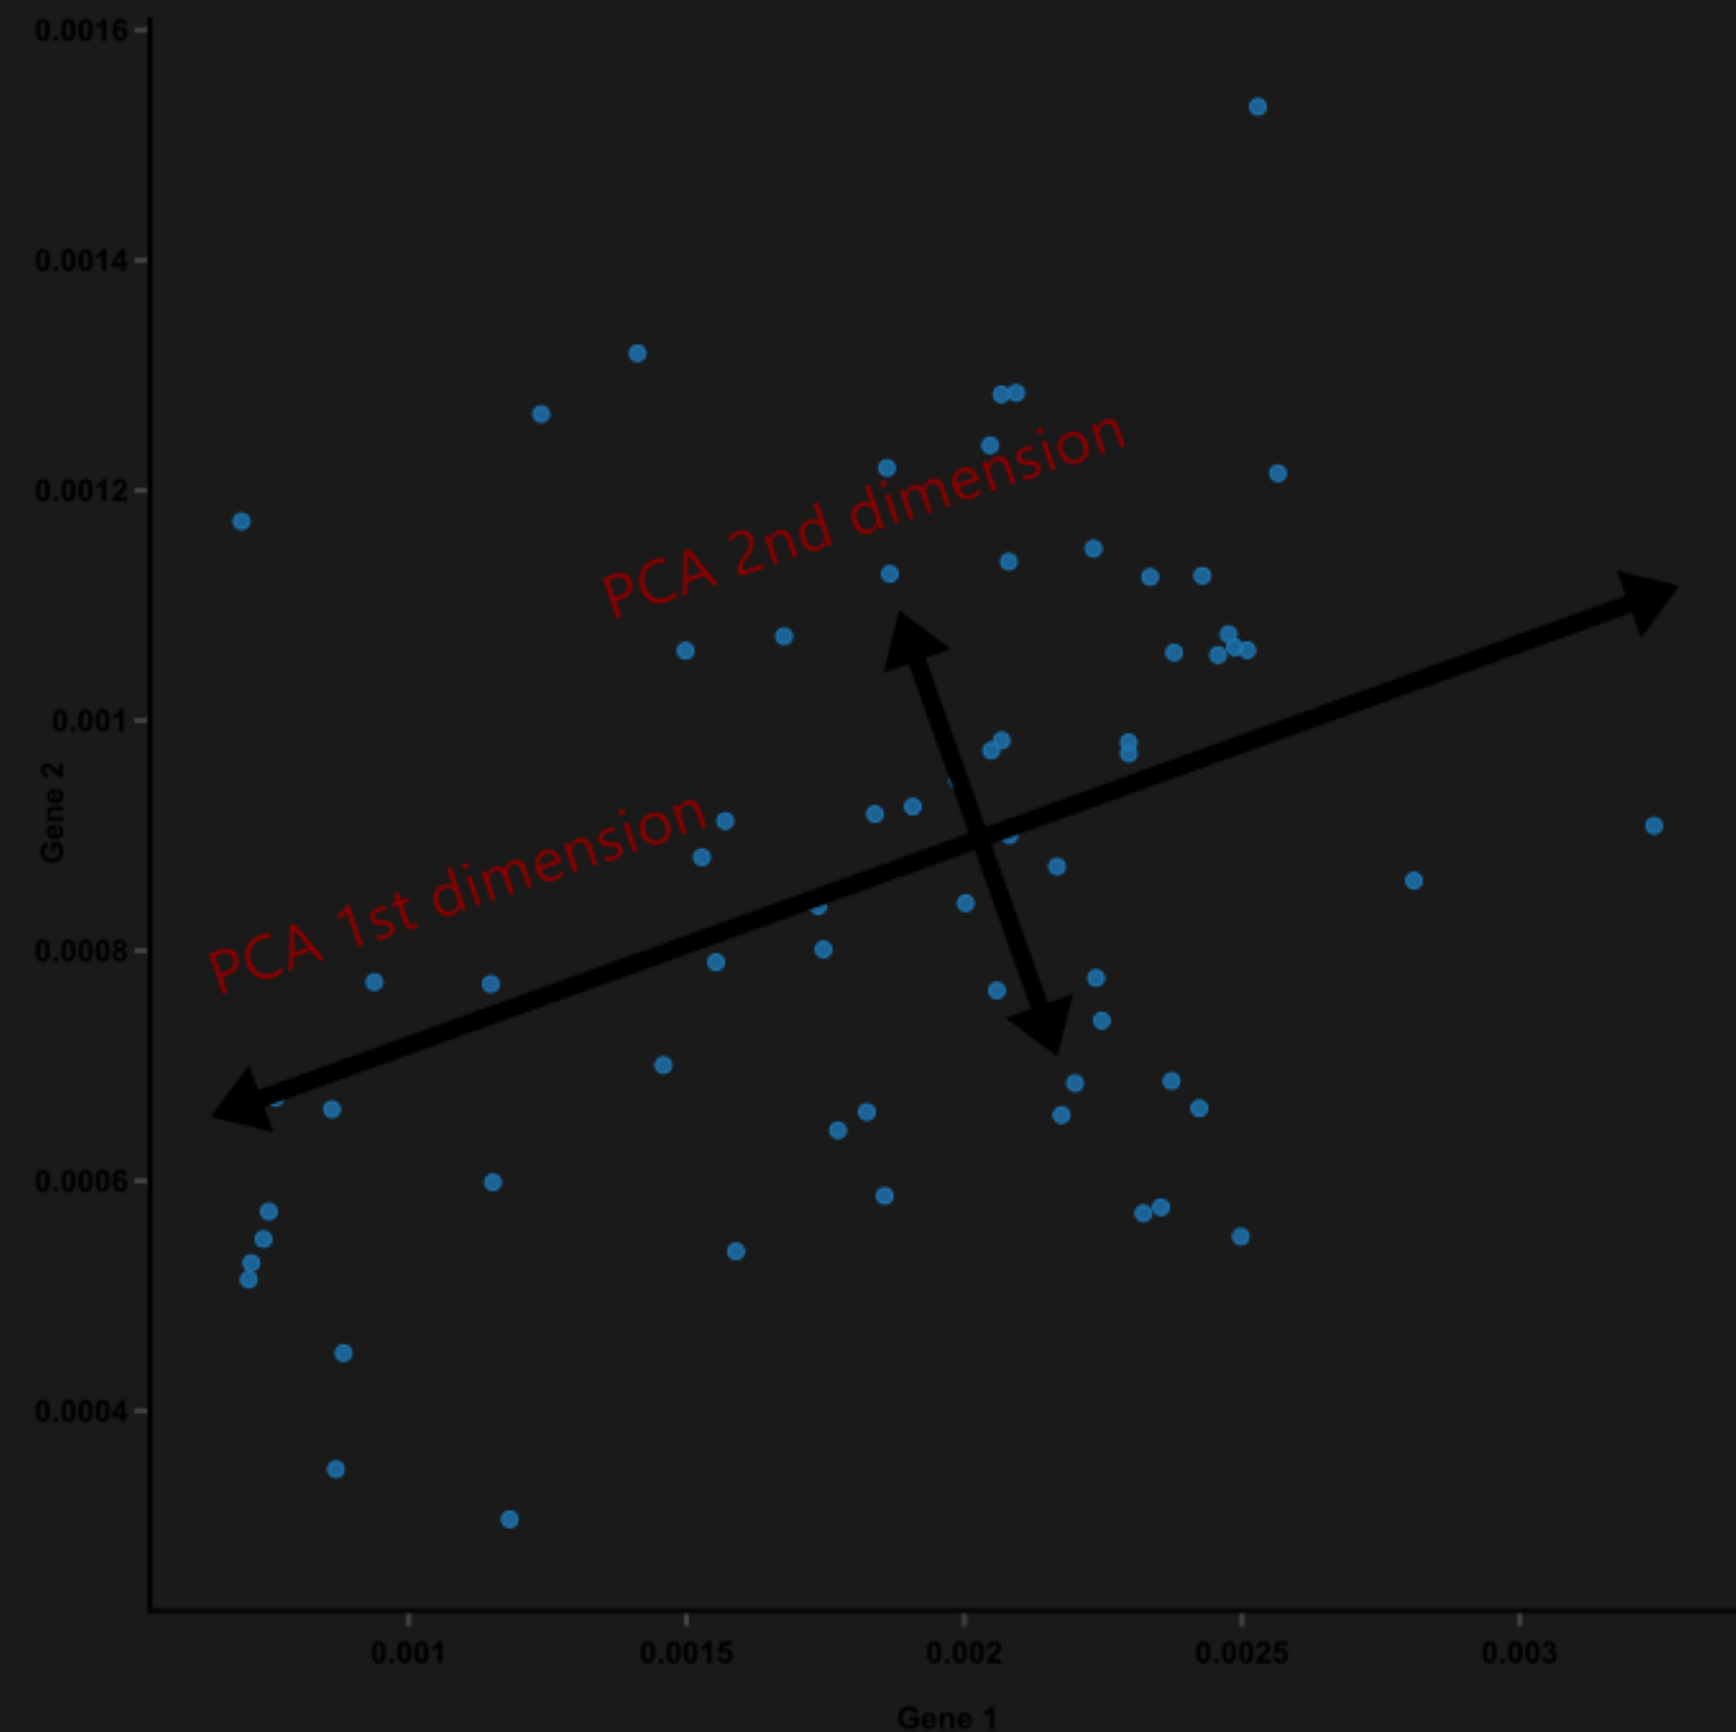
As we can see the
output of first
two columns is
very similar.

	PValueSelector	JensenShannonSelector	BregmanDivergenceSelector
0	age	age	workclass_0
1	workclass_2	workclass_2	workclass_1
2	workclass_3	fnlwgt	workclass_2
3	fnlwgt	education_2	workclass_3
4	education_1	education_3	education_0
5	education_2	education_4	education_1
6	education_3	education.num	education_2
7	education_4	marital.status_0	education_3
8	education.num	marital.status_2	education_4
9	marital.status_0	occupation_1	marital.status_0
10	marital.status_1	occupation_2	marital.status_1
11	marital.status_2	relationship_0	marital.status_2
12	occupation_0	relationship_2	occupation_0
13	occupation_1	race_2	occupation_1
14	occupation_2	sex	occupation_2
15	occupation_3	capital.gain	occupation_3
16	relationship_0	hours.per.week	relationship_0

Data Correlation



Principal component analysis & ICAReducer



PCA from 36 to 30 columns

ICAReducer from 36 to 36

Kydavra **ICAReducer**
for reducing
the dimensionality
of your data



Feature Scaling - Standardization

	0	1	2	3	4	5	6	7	8	9	...	26	27	28	29
0	0.872712	-0.019407	-0.442500	0.510810	-0.462514	0.092040	-0.032815	-0.424906	-0.547292	1.048787	...	-0.196578	-0.255178	0.267353	-0.08584
1	-1.470061	-0.019407	-0.442500	0.510810	-0.462514	-0.795088	-0.032815	-0.424906	1.827179	-0.953482	...	-0.196578	-0.255178	-1.042844	-0.08584
2	1.195853	-0.019407	-0.442500	0.510810	-0.462514	2.608657	-0.032815	-0.424906	1.827179	1.048787	...	-0.196578	-0.255178	1.577549	-0.08584
3	-0.419853	-0.019407	-0.442500	0.510810	-0.462514	-0.450815	-0.032815	2.353460	-0.547292	1.048787	...	-0.196578	-0.255178	-0.518765	-0.08584
4	-0.339067	-0.019407	-0.442500	0.510810	-0.462514	0.391721	-0.032815	-0.424906	-0.547292	-0.953482	...	-0.196578	-0.255178	-0.606112	-0.08584
...
37181	-0.500638	-0.019407	-0.442500	0.510810	-0.462514	-1.066124	-0.032815	-0.424906	-0.547292	-0.953482	...	-0.196578	-0.255178	1.053471	-0.08584
37182	0.791926	-0.019407	-0.442500	0.510810	-0.462514	-0.295012	-0.032815	-0.424906	-0.547292	-0.953482	...	-0.196578	4.373423	-0.169380	-0.08584
37183	0.388000	-0.019407	-0.442500	0.510810	-0.462514	1.079886	-0.032815	-0.424906	1.827179	1.048787	...	-0.196578	-0.255178	0.704085	-0.08584
37184	0.064859	-0.019407	-0.442500	0.510810	-0.462514	-0.750734	-0.032815	-0.424906	1.827179	1.048787	...	-0.196578	-0.255178	0.704085	-0.08584
37185	0.630356	-0.019407	2.259885	-1.957674	2.162095	-0.719079	-0.032815	-0.424906	-0.547292	1.048787	...	-0.196578	-0.255178	-0.169380	-0.08584
37186 rows × 36 columns															

Standardization is less sensitive to outliers compared to normalization (Min-Max scaling).


Model Training

-SVC

-RandomForestClassifier

-MLPClassifier

-KNeighborsClassifier



Model: SVC

```
elif isinstance(model, SVC):  
    clf = SVC()  
    clf.fit(X_train, y_train)  
    output69 = clf.predict(X_test)  
    f1 = accuracy_score(y_test, output69)  
    lista4.append(f1)
```

```
fun1(SVC())
```

```
Simple df  
0.8911108669376991  
PCA dataframe  
0.8910009889023184  
Isolation Forest dataframe  
0.8929225137278829  
EllipticEnvelope dataframe  
0.8860944939567819  
LocalOutlierFactor dataframe  
0.8881423856671381
```

Model: RandomForestClassifier

```
elif isinstance(model, RandomForestClassifier):
    tuned_parameters = {
        'criterion': ['gini', 'entropy', 'log_loss'],
        'max_features': ['log2', 'sqrt', None],
        'class_weight': ['balanced', 'balanced_subsample']
    }
    clf = GridSearchCV(RandomForestClassifier(), tuned_parameters, scoring='recall')
    clf.fit(X_train, y_train)
    output69 = clf.predict(X_test)
    f1 = accuracy_score(y_test, output69)
    lista2.append(f1)
```

```
fun1(RandomForestClassifier())
```

```
Simple df
0.8978134270959235
PCA dataframe
0.8905614767607956
Isolation Forest dataframe
0.8942953020134228
EllipticEnvelope dataframe
0.8902453912831156
LocalOutlierFactor dataframe
0.8892032060348892
```

Model: KNeighborsClassifier

```
elif isinstance(model, KNeighborsClassifier):
    tuned_parameters = {
        'weights': ['uniform', 'distance'],
        'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute']
    }
    clf = GridSearchCV(KNeighborsClassifier(), tuned_parameters, scoring='recall')
    clf.fit(X_train, y_train)
    output69 = clf.predict(X_test)
    f1 = accuracy_score(y_test, output69)
    lista3.append(f1)
```

```
fun1(KNeighborsClassifier())
```

Simple df

0.8674870893308427

PCA dataframe

0.8692451378969344

Isolation Forest dataframe

0.87614399023795

EllipticEnvelope dataframe

0.8670492003418386

LocalOutlierFactor dataframe

0.8640971239981141

Model: MLPClassifier

```
elif isinstance(model, MLPClassifier):
    tuned_parameters = {
        'activation':['identity', 'logistic', 'tanh', 'relu'],
        'solver': ['lbfgs', 'sgd', 'adam'],
        'learning_rate':['constant', 'invscaling', 'adaptive']
    }
    clf = GridSearchCV(MLPClassifier(), tuned_parameters,scoring='recall')
    clf.fit(X_train, y_train)
    output69 = clf.predict(X_test)
    f1 = accuracy_score(y_test,output69)
    lista8.append(f1)
```

```
fun1(MLPClassifier())
```

Simple df

0.878694648939677

PCA dataframe

0.8845181848148556

Isolation Forest dataframe



0.5846552776082977

EllipticEnvelope dataframe

0.8879257721889879

LocalOutlierFactor dataframe

0.8843705799151343





Model: Hyperparameter tuning

```
from sklearn.model_selection import GridSearchCV
tuned_parameters = {
    'criterion': ['gini', 'entropy', 'log_loss'],
    'max_features': ['log2', 'sqrt', None],
    'class_weight': ['balanced', 'balanced_subsample']
}
clf = GridSearchCV(RandomForestClassifier(), tuned_parameters, scoring='recall')
clf.fit(X_train, y_train)
output69 = clf.predict(X_test)
f1 = accuracy_score(y_test, output69)
f1
```

▼	RandomForestClassifier
	RandomForestClassifier(class_weight='balanced_subsample', criterion='entropy', max_features=None)

Model Interpretation

RandomForestClassifier



Mulumesc pentru Atentie!